

Supervisory earnings pipeline

The question

Not whether a call can be summarised. Whether earnings Q&A carries a **leading** signal that reported metrics **alone** do not.

1,500

Firms the PRA supervises

137

UK results transcripts

0.28

FinBERT kappa, not 60%

WATCH

HSBC H1 2025

These slides are the story told in the pitch, closed with the cases held for this report. The notebook is the technical walkthrough.

Structured returns are the pipeline. The call is not.

What the PRA already runs

About 1,500 firms. Returns on a defined schedule, checked against defined fields.

What stays outside

Quarterly results: a curated presentation plus unscripted Q&A.
Expensive to read properly, so the signal is not used systematically.

What we measure besides tone: was the answer direct, did it supply the metric, or did it change the subject. Positive and evasive can be the same sentence.

Two words

Leading: did the call move before the stress, not after. Alone: information the four reported lines do not already contain.

Positive tone, evasive answer

Analyst: the 90-day past-due rate on the UK cards book, this quarter against the last. Management: comfortable with the overall credit picture; the portfolio remains robust.

What a tone model writes	What we also write
Positive. The reply sounds confident.	Directness low. The number asked for was not given.
One label for the whole answer.	Topic substituted: UK cards became “overall credit”.

Why it matters
If directness on one topic falls while the tone stays cheerful, across 62 paired quarters, that is the warning. Sentiment alone is not enough.

Same chain as the pitch. Now it is the product.

Step	What happens	Owner
1 Ingest	PDFs and Excel. Files stay inputs.	Susana
2 Prepare	Speaker split, then analyst Q&A pairs.	Susana
3 Model	BERTopic, FinBERT, lexicon, behaviour, eight-way map.	Debanjan · Alfred · Bupathi · Aidan
4 Compare	Time, matched peer, pack versus Q&A.	Susana · Aidan
5 Decide	Rules to alert, watch, or null.	Aidan · Bupathi
6 Publish	PRA note. Notebook is the factory.	Rafael · Susana

Re-run

45–90 minutes when new IR files land. Recalibration is a human gate, not a job that retrains itself.

137 files. Gaps left as gaps.

Fact from the build	Figure
UK transcripts. One equity-analyst meeting excluded.	137 files, 136 results calls
Quarters where both banks published. Paired by quarter, not by full year.	62 quarters
Speaker turns, then Q&A pairs. 192 replies sat inside an analyst turn.	4,188 turns · 1,213 pairs
Answers that address the four metrics, among questions about them.	600 of 818 · 73%
Empty answers, and rows with no publication date. Kept, not patched.	108 · 369

HSBC is twice a year until 2008. Barclays starts in 2010 Q2. Nothing is interpolated or filled with zero. Credit Suisse was built and held out of the pitch. It is in this report, and it is not in the peer gap.

Categories first. Then the models.

We kept	Why this order
18 PRA categories, before any transcript	A map built from the calls would succeed by design. Liquidity barely mentioned is a finding.
BERTopic, LDA as the check	Coherence 0.46 versus 0.40. Five topics. Ids are not permanent alert keys.
Whole-turn FinBERT	Lexicon kappa 0.02, chance. Sentence scoring kappa 0.13. Neither replaces the whole answer.
Behaviour, not a fourth sentiment class	Direct, metric supplied, topic substituted. Avoidance is behaviour.

Fine-tune stays off

A gain on labels the model helped to create is not a promotion. The second coder is still the gate. We promote only if the fine-tune beats plain FinBERT on human labels.

The 70% bar was not met

0.28

FinBERT kappa

0.02

Lexicon kappa. Chance.

80%

Questions called neutral

70%

Bar set first. Not met.

Say kappa, not 60%

Both methods reach for neutral, so raw agreement flatters them.
FinBERT does something. The lexicon does not. They agree with each other about 58% of the time.

What we did not move

The 70% bar was set before coding. We published the miss.
Directness is triage: where to look, not a judgement of a firm or a person.

Anachronism

The PRA did not exist before 2013. A category from today's PRA sources, laid on a 2006 transcript, is anachronistic. The run is provisional for that reason.

HSBC 2025 is watch. The peer gap is not an alert.

-0.13

FinBERT net

8

Analyst turns

n/a

Metric hits in the call

WATCH

Rule A3

Why alert did not fire

Matched Barclays on the same half-year is entirely FinBERT-neutral (n=7). The gap is about -0.13. Rule A2 refuses it. Absence of a metric hit is not recorded as agreement.

The pack shows the quote and the sample size. It does not rate HSBC.

The other two reviewed cases

Case	Verdict
Barclays 2026-H1, net -0.11, n=7	Watch. Peer refused.
Credit Suisse 2022-Q4, net +0.08, n=20	Null. Not a predicted collapse.

What the finished run shows

Finding	Number	Do not claim
Q&A is not a copy of the pack	Agree on 26% of 212 overlaps	That disagreement is automatically an alert
Credit narrative is the weak join	18% agree, 44 rows	That ECL is mis-stated
Soft, without a bad peer, is watch	HSBC -0.13, rule A3	That the firm is weaker
A known failure can score null	CS Q4 2022, FinBERT +0.08	That the pipeline saw the collapse coming

BERTopic coherence 0.46, LDA 0.40. Both stay. Only BERTopic names the live themes, and even those names are not used as permanent alert keys.

What the Bank should do with this

Use now

Run the pack when filings land. A person accepts or rejects watch, alert, or null. Keep the peer-quality gate. File a null in the same way as a watch.

Do not do

Do not promote the fine-tuned head. Do not add banks to make the peer gap look busier. Do not treat neutral as “no issue”.

Next analysis

Second coder on the 90-pair gold. Retag the four metrics inside the episode so the agree column is evidenced rather than n/a. Only then reopen promotion.

Limit

Eight turns is a print, not a population. Hedged answers still collapse toward neutral. That limit is on the note.

These slides are the story. The notebook is the proof.

Open the notebook, not every cell

Stage	Show
0	Colab clones the repo, then finds scripts/
1	PDFs to speaker turns to Q&A pairs
2–3	BERTopic, then FinBERT on question and answer
6	Three baselines, each with n
8–9	Rules, then the PRA note

For a technical listener

State lives in data/boe.sqlite. CSV copies are in data/processed/.
Human labels are scored, not rewritten.

If a live re-fit is too slow

Walk the saved outputs already in the notebook, then the PRA note for HSBC 2025-H1. Do not refit BERTopic on stage.

Sentiment alone is not enough

Close

The strongest result is a negative one. FinBERT and the lexicon agree about 58% of the time. FinBERT calls about 80% of questions neutral and misses the hedged negatives. Kappa is 0.28. We did not move the 70% bar. HSBC half-year 2025 is watch, not a rating. Credit Suisse's failure quarter, held out of the pitch, is null.

For the Bank

Is alert / watch / null the right shape, and is a 45–90 minute re-run acceptable each quarter?

For markers

Depth on two banks, an honest miss on the 70% bar, and a pipeline that can say null.